Goedel-Prover: A Frontier Model for Open-Source Automated Theorem Proving

Yong Lin[*†] **Shange Tang**[*†] **Bohan Lyu**[‡] **Jiayun Wu**[‡] **Hongzhou Lin**[§] **Kaiyu Yang**[¶]

Jia Li[‖] **Mengzhou Xia**[†] **Danqi Chen**[†] **Sanjeev Arora**[†] **Chi Jin**[†]

Abstract

We introduce Goedel-Prover, an open-source language model that achieves state-of-the-art (as of April 5 2025) performance in automated formal proof generation for mathematical problems. A key challenge in this field is the scarcity of formalized mathematical statements and proofs, which we address through the following approaches. First, we train LLMs to convert natural language math problems from the Numina dataset to equivalent formal statements in Lean 4. This process creates the dataset Goedel-Pset-v1, which includes 1.64 million formal statements. Next, we develop a large dataset of formal proofs by training a series of provers. Each new prover can prove many statements that previous ones could not, and these new proofs are added to the training set for the next prover. Finally, we obtain the dataset Goedel-Pset-v1-solved, which contains proofs for over 800K statements from Goedel-Pset-v1. Supervised fine-tuning (SFT) of DeepSeek-Prover-V1.5-Base on Goedel-Pset-v1-solved (i.e., no RL) yields a model that achieves a success rate of 57.6% (Pass@32) on miniF2F benchmark, surpassing the previous leader DeepSeek-Prover-V1.5 (trained using SFT + RL on a proprietary dataset) by 7.6%. On PutnamBench, Goedel-Prover-SFT successfully solves 7 problems (Pass@512), ranking first on the leaderboard. Further RL training (including DPO) improves Goedel-Prover-SFT's success rate to over 60% (Pass@32) on miniF2F.

To aid future research, we provide extensive discussion of our training methodology and design choices. We also fully open-source our codes, models, and datasets. Additionally, we open-source formal proofs for 29.7K problems in Lean Workbook, nearly doubling the 15.7K solved by prior provers.

1 Introduction

Recent advancements in large language models (LLMs) have demonstrated remarkable capabilities in reasoning tasks, especially in solving mathematical problems (Guo et al., 2025; Yang et al., 2024a). These models excel at reasoning through natural language, which we refer to as *informal reasoning*. However, natural language-based reasoning is difficult to automatically verify by machines, which undermines the reliability of informal reasoning in practical applications. This also makes it more difficult to further improve the reasoning capabilities of language models. In contrast to informal reasoning, *formal reasoning* allows reasoning in a machine-verifiable format, opening up new possibilities for verification and automation. In particular, proof assistants such as Lean (De Moura et al., 2015; Moura & Ullrich, 2021), Isabelle (Paulson, 1994), and Coq (Barras et al., 1997) provide formal

---

[*]YL and ST contribute equally to this work. {yong.lin,shangetang}@princeton.edu

[†]Princeton Language and Intelligence, Princeton University.

[‡]Tsinghua University.

[§]Amazon. This work is independent of and outside of the work at Amazon.

[¶]Meta FAIR. KY served an advisory role. All experiments were conducted on PLI servers.
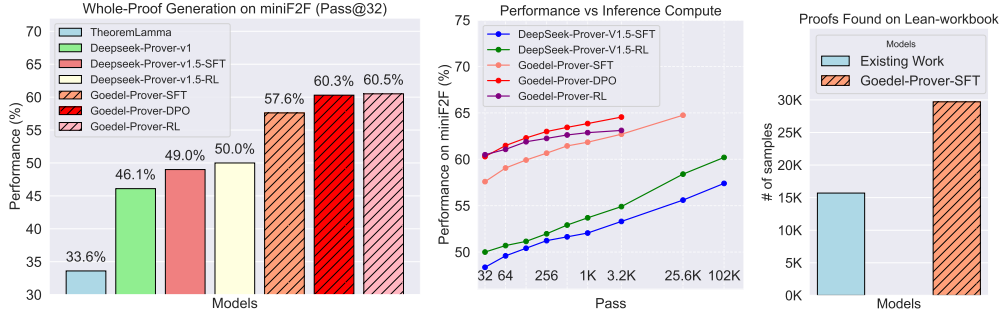
[‖]Numina.

Figure 1: (**Left**) Pass@32 performance on miniF2F for whole-proof generation, compared to previous SOTA models. (**Middle**) A comparison of Goedel-Proverand DeepSeek-Prover-V1.5 on miniF2F performance across varying inference budgets, ranging from Pass@32, 64, 128, ..., to 4 × 6400. (**Right**) Goedel-Prover-SFT solves 29.7K problems in the Lean Workbook. In comparison, InternLM2.5-Step-Prover (Wu et al., 2024) and InternLM-Math-Plus (Ying et al., 2024b) collectively solved 15.7K samples.

languages that can express reasoning in a way that can be mechanically verified. Thus, it is of great interest to train LLMs to write proofs in these formal languages.

A significant challenge in training LLMs for theorem proving in formal languages is the scarcity of formalized math statements and proofs. Writing proofs for theorems expressed in formal languages is highly demanding and necessitates considerable domain expertise. Therefore, existing publicly available datasets for formal languages are limited in size. For example, the Lean Workbook (including Lean Workbook Plus) dataset (Ying et al., 2024a; Wu et al., 2024) comprises a total of 140K formal statements, where formal statements refer to problem statements in Lean without proofs. However, only 15.7K of these statements come with formal proofs, which were found by InternLM2.5-StepProver and InternLM-Math-Plus (Ying et al., 2024a; Wu et al., 2024; Ying et al., 2024b). Additionally, the Open Bootstrapped Theorems dataset (Wang et al., 2024) includes 107K statements with proofs sourced from Mathlib4 (mathlib4, 2023). However, Mathlib4 exhibits significant distribution shift from general problem-solving benchmarks, such as the widely used miniF2F (Zheng et al., 2021). See Section 5 for detail.

In contrast to the scarcity of data in formal languages, there is a vast amount of math problems and solutions written in informal language. For example, Numina (Li et al., 2024a) includes 860K high-quality question and answer pairs sourced from MATH (Hendrycks et al., 2021), GSM8K (Cobbe et al., 2021), AMC (aop), AIME (MAA, 2024), the AoPS Forum (aop), Chinese K-12 Exams (Shao et al., 2024), World Olympiads, and synthetic data (Mitra et al., 2024). We start by training LLMs to formalize the problem statements in this dataset into Lean. To increase the diversity of the formalization styles, we train two formalizers. One is trained on informal and formal (I-F) statement pairs from Lean Workbook, while the other is trained on I-F statement pairs annotated by Claude-sonnet-3.5 (Anthropic, 2024). We use these two formalizers to formalize the statements and then employ LLMs to ensure that the formal statements preserve the content of the informal statements. Our efforts result in 1.64 million formal statements.

Using this extensive dataset of formal statements, we employ expert iteration (Polu et al., 2022) to train the prover to generate proofs. Notably, we train a model to generate complete proofs solely based on statements, without interacting with the Lean compiler during the generation process. This approach is referred to as the whole-proof generation method (Jiang et al., 2022; Wang et al., 2024; Xin et al., 2024a;b). At the beginning of the expert iteration, we generate 16 proof candidates using DeepSeek-Prover-V1.5-RL (the previous SOTA) for each formal statement in our dataset, and then we verify the correctness of each candidate using Lean compiler. The correct proofs are then collected to train our iter-1 prover based on DeepSeek-Prover-V1.5-Base. In subsequent rounds, we utilize our iter-k prover to collect new proofs and add them to the training data. We then perform supervised fine-tuning

starting from DeepSeek-Prover-V1.5-Base for another round, resulting in the iter-$(k + 1)$ prover. We conduct a total of 8 iterations and observe a consistent improvement starting from the first iteration.

We demonstrate that expert iteration, with large-scale formalized statements can lead to SOTA performance in formal proof generation. Specifically,

- Our model, Goedel-Prover-SFT, outperforms DeepSeek-Prover-V1.5-RL (the previous SOTA model) by 7.6% on miniF2F, achieving a Pass@32 score of 57.6% (i.e., the prover generated 32 proofs for a problem, and 57.6% of the problems have at least one correct proof verified by Lean) compared to DeepSeek-Prover-V1.5-RL's 50.0%, as shown in Figure 1 (left). It consistently surpasses DeepSeek-Prover-V1.5-RL across all sampling budgets, including Pass@32, 64, and up to 25600, as shown in Figure 1 (middle).

- We have cumulatively solved 29.7K problems in Lean Workbook, significantly increasing the existing 15.7K proofs found by InternLM2.5-StepProver and InternLM-Math-Plus (Wu et al., 2024; Ying et al., 2024b), as shown in Figure 1 (right).

- Goedel-Prover-SFT solves 7 problems on PutnamBench by Pass@512[1], securing the **#1** position on the leaderboard (Table 2).

- We open source our codes[2], models[3 4 5 6], datasets[7 8], and the new proofs discovered[9] in the Lean Workbook to facilitate future research.

To understand the factors behind Goedel-Prover's strong performance, we provide an in-depth discussion of our training recipe, analyzing the effects of scaling up training data, the diversity introduced by autoformalization, correlations among datasets, and alternative data synthesis strategies. Furthermore, although our final model is trained purely through supervised fine-tuning, we also explore direct preference optimization (DPO) and reinforcement learning (RL) techniques built on top of it. Our Goedel-Prover-DPO and Goedel-Prover-RL achieves a success rate over 60% (Pass@32) on miniF2F. However, we also find that DPO and RL-trained models tend to overfit to "shortcuts" and benefit less from increased inference-time compute.

2 Related Work

Automated theorem proving. Automated theorem proving (ATP) is a long-standing problem in symbolic AI (Robinson & Voronkov, 2001). Traditional approaches represent theorems in first-order logic and prove them using decision procedures (De Moura & Bjørner, 2008; Barbosa et al., 2022) and search (Kovács & Voronkov, 2013; Schulz et al., 2019). The proof search has been enhanced by replacing handcrafted heuristics with machine learning techniques (Urban et al., 2011; Kaliszyk et al., 2018). However, approaches based on first-order logic struggle to scale to complex theorems and often do not yield human-readable proofs.

In recent years, learning-based theorem proving has undergone a significant transformation. A notable approach, introduced by Polu & Sutskever (2020), involves leveraging large language models to assist in theorem proving with proof assistants such as Lean (De Moura et al., 2015; Moura & Ullrich, 2021) and Isabelle (Paulson, 1994). Follow-up research has explored various avenues, such as retrieving useful lemmas (Irving et al., 2016; Mikuła et al.,

---

[1]We initially solved 8 problems on PutnamBench. However, after discussing with the authors of PutnamBench, we discovered that one of the problems was mis-formalized. Therefore, this problem is not included in our count, and we report a total of 7 problems here.

[2]https://github.com/Goedel-LM/Goedel-Prover

[3]https://huggingface.co/Goedel-LM/Goedel-Prover-SFT

[4]https://huggingface.co/Goedel-LM/Goedel-Prover-DPO

[5]https://huggingface.co/Goedel-LM/Goedel-Formalizer-32B-SonnetAnnotated

[6]https://huggingface.co/Goedel-LM/Goedel-Formalizer-32B-LeanWorkbookAnnotated

[7]https://huggingface.co/datasets/Goedel-LM/Goedel-Pset-v1

[8]https://huggingface.co/datasets/Goedel-LM/Goedel-Pset-v1-solved

[9]https://huggingface.co/datasets/Goedel-LM/Lean-workbook-proofs

2024; Yang et al., 2024b), utilizing Monte Carlo tree search for proof discovery (Lample et al., 2022), and harnessing the capabilities of large language models (LLMs) for natural language reasoning (Jiang et al., 2022; Lin et al., 2024). Notably, Polu et al. (2023) was the first to apply expert iteration (Anthony et al., 2017) to theorem proving. This method alternates between two phases: (1) attempting to prove unsolved theorems and (2) enhancing the prover by incorporating newly discovered proofs into its training data. Expert iteration has yielded significant improvements in several recent provers (Wu et al., 2024; Xin et al., 2024b), including our Goedel-Prover.

Most automated theorem provers operate in a stepwise manner, generating individual proof steps that are then assembled into complete proofs using proof search algorithms. Recently, researchers have shown that generating entire proofs is feasible (First et al., 2023; Xin et al., 2024a; Wang et al., 2024). This approach avoids the costly search process, resulting in lower latency and potentially offering a more efficient use of computational resources during testing. While Goedel-Prover also generates whole proofs, our data and methodology can, in principle, be adapted to develop stepwise provers as well.

Autoformalization and synthetic data generation. The shortage of high-quality formal mathematical data poses a significant bottleneck in training theorem-proving models. While techniques like reinforcement learning may reduce the reliance on human-written proofs (Google DeepMind, 2024), there remains a need for a substantial number of formal theorem statements. A promising approach is to synthesize formal statements through autoformalization, where large language models (LLMs) translate informal mathematical statements into formal ones (Wu et al., 2022; 2024; Xin et al., 2024a;b).

DeepSeek-Prover (Xin et al., 2024a) and InternLM2.5-StepProver (Wu et al., 2024) have successfully implemented this strategy to formalize a large volume of statements into Lean for expert iteration. We adopt a similar approach. The difference is: while Liu et al. (2024) focuses on formalizing their internal dataset, we concentrate on formalizing the Numina dataset (Li et al., 2024a) alongside a privately collected dataset. Additionally, we train two formalizers to enhance the diversity of formalization styles, which we demonstrate to be beneficial in Section 4.

3 Method

We begin by translating informal statements (expressed in natural language) into formal statements (represented in Lean). Using these formal statements, we iteratively train our prover with proofs generated by the prover and verified by the Lean compiler. The details of each step are elaborated in the following parts.

3.1 Statement Formalization

We first train the statement formalizers to translate informal statements in Numina into formal statements as shown in Figure 2. To enhance the diversity of formalized statements, we train two models to formalize informal statements.

- **Formalizer A:** We train the Formalizer A model using Formal and Informal (F-I) statement pairs sourced from Lean Workbook.
- **Formalizer B:** We employ Claude-sonnet-3.5 to formalize 230K statements from Numina. From this set, we extract 170K statements that successfully passed Lean compilation. These 170K F-I statement pairs are then used to train Formalizer B.

Both Formalizer A and B are trained using supervised fine-tuning with Qwen2.5-Coder-32B[10]. The training of these two formalizers takes less than 24 hours on 8 H100 GPUs. See Appendix A.1 for examples of formalized statements form two formalizers, where we observe that even for the same problem, the style of the formalized statement can impact the prover's performance.

---

[10]https://huggingface.co/Qwen/Qwen2.5-Coder-32B

Formal and Informal (F-I) statement pair

| Informal statement	Let a, b, and c be real numbers that satisfying $3a+b+c=-3$, $a+3b+c=9$, and $a+b+3c=19$. Prove that the product $abc=-56$.
Formal statement	`theorem` mathd_algebra_338 `(a b c : R)` (h₀ : 3 * a + b + c = -3) (h₁ : a + 3 * b + c = 9) (h₂ : a + b + 3 * c = 19) : a * b * c = -56 := `by`

Figure 2: This figure illustrates the training of the formalizers. The term "F-I statement pairs" refers to pairs consisting of Formal and Informal (F-I) statements. An example is shown on the right part. We train two formalizers, Formalizer A and B, using F-I statement pairs sourced from various origins.
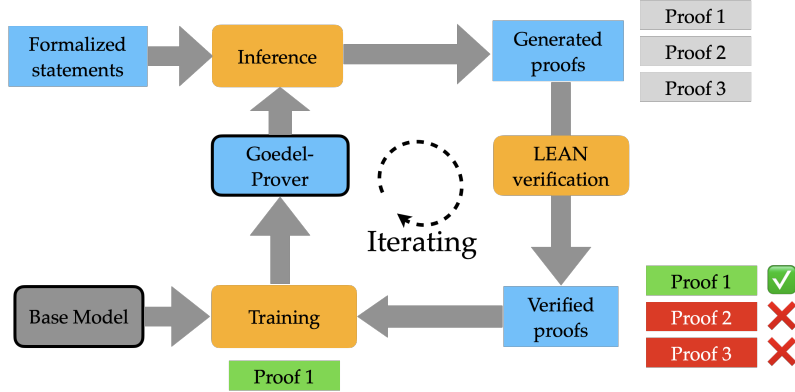

Figure 3: This figure illustrates the process of expert iteration. Each time, we utilize our iter-k prover to collect new proofs and add them to the training data. We then conduct supervised fine-tuning starting from DeepSeek-Prover-V1.5-Base for another round, resulting in the iter-$(k+1)$ prover.

Quality assessment. We employ two tests to assess the quality of the formalized statements. First, the formalized statement must conform to Lean syntax and can successfully compile, with the potential proof replaced by the placeholder ":= by sorry". This syntax check is known as the Compiling Correctness (CC) Test in the literature (Ying et al., 2024a). Second, the formalized statement must accurately capture the original informal problem, incorporating all assumptions, conditions, and implicit definitions. We refer to this second test as the Faithfulness and Completeness (FC) Test. For the FC test, we use Qwen2.5-72B-Instruct[11], details are presented in Appendix A.2.

In addition to formalizing the 860K open-sourced Numina (Li et al., 2024a) datasets, we also formalize a private 68K collection of math problems from Art of Problem Solving (AOPS), which has been collected and processed by the Numina group (Li et al., 2024a). Out of a total of 928K informal statements, 760K have two valid formalized statements generated by Formalizer A and B, while 123K contain only one valid formalized statement. After formalizing both the Numina and AOPS datasets, we further incorporate 140K statements from Lean Workbook, including Lean Workbook Plus. As a result, we have a total of 1.78M formal statements.

3.2 Expert Iteration

After obtaining a large collection of formalized statements in Section 3.1, we employ expert iteration to train the prover (Liu et al., 2024; Wu et al., 2024; Li et al., 2024b), which is

---

[11] https://huggingface.co/Qwen/Qwen2.5-72B-Instruct

illustrated in Figure 3. Specifically, we first utilize DeepSeek-Prover-V1.5-RL[12] to generate 16 proofs for each statement. We then verify these proofs with the Lean compiler. If at least one proof solves the statement, we retain one proof per statement. In cases where multiple proofs are available, we randomly sample one solution. These collected proofs are used for supervised fine-tuning (SFT) based on DeepSeek-Prover-V1.5-Base[13], resulting in the iter-1 prover. We continue this expert iteration process; each time, we use the iter-k prover to generate answers and cumulatively collect correct solutions to train DeepSeek-Prover-V1.5-Base for the next iteration, the iter-$(k+1)$ prover. Refer to Appendix B for more details on each iteration.

We experiment with learning rates of 1×10^{-4} and 5×10^{-5}, training for either 1 or 2 epochs. We use the packing trick (Tunstall et al., 2022) with a small batch size of 8 to speed up the training. In each iteration, the training time for 1 epoch is approximately 12 hours using 4 H100 GPUs. The inference time for the 1.78M statements set by Pass@16 is 6 hours, utilizing 64 H100 GPUs. Additionally, the verification time for these proofs requires 10 hours with 8,000 CPUs.

4 Results

Benchmarks. Following the works of (Wang et al., 2024; Xin et al., 2024a; Wu et al., 2024; Li et al., 2024b), we primarily use miniF2F (Zheng et al., 2021) as our main evaluation benchmark. We also track the problems solved by our prover in Lean Workbook (Ying et al., 2024a) and investigate the performance on ProofNet (Azerbayev et al., 2023) and PutnamBench (Tsoukalas et al., 2024). Additionally, we uniformly sample a subset from our formalized dataset to create a held-out evaluation dataset. Below, we provide descriptions of each dataset.

- miniF2F (Zheng et al., 2021) is a formal theorem proving benchmark, consisting of 488 problem statements (244 validation and 244 test problems) in Lean. The problems are drawn from high-school exercises, as well as high-school level competitions including the AIME, AMC, and the International Mathematical Olympiad (IMO). The original benchmark was released in Lean 3, and for our analysis, we use the version of miniF2F in Lean 4.9.0 provided by Xin et al. (2024a).
- ProofNet (Azerbayev et al., 2023) is a formal theorem proving benchmark of undergraduate-level mathematics, consisting of 371 problem statements in Lean (185 validation and 186 test problems). The problems are primarily drawn from undergraduate pure mathematics textbooks, covering topics such as real and complex analysis, linear algebra, abstract algebra, and topology. The original benchmark was released in Lean 3, and for our analysis, we use the version of ProofNet in Lean 4.9.0 provided by Xin et al. (2024a).
- Lean Workbook (Ying et al., 2024a) is a large-scale Lean 4 problem set formalized from natural language math problems (mainly from the forum AOPS), which consists of 140K statements in Lean 4. We also monitor the problems solved by our model during the expert iteration process. Notably, the problem set from Lean Workbook is included in this training, which is consistent with DeepSeek-Prover-V1.5 (Xin et al., 2024a) and InternLM2.5-StepProver (Wu et al., 2024).
- PutnamBench (Tsoukalas et al., 2024) is a formal theorem proving benchmark on competition mathematics problems sourced from the William Lowell Putnam Mathematical Competition years 1962 - 2023. PutnamBenchcomprises 644 Lean 4 statements, covering algebra, analysis, number theory, geometry, combinatorics, probability and set theory.
- NuminaTest. We randomly sample 250 statements from our formalized Numina dataset and use it as a held-out testing set. We refer to this subset as NuminaTest.

Main results. The performance on miniF2F is shown in Table 1. The Pass@32 performance of our Goedel-Prover-SFT is 57.6%, surpassing the previous SOTA open source model,

---

[12]https://huggingface.co/deepseek-ai/DeepSeek-Prover-V1.5-RL
[13]https://huggingface.co/deepseek-ai/DeepSeek-Prover-V1.5-Base

Whole-Proof Generation Model	Pass	Performance
TheoremLamma (Wang et al., 2024)	128	33.6%
Deepseek-Prover-V1 (Xin et al., 2024a)	32	46.1% ± 0.5%
DeepSeek-Prover-V1.5-SFT (Xin et al., 2024b)	32	48.2% ± 0.6%
DeepSeek-Prover-V1.5-RL (Xin et al., 2024b)	32	50.0% ± 0.5%
Goedel-Prover-SFT	32	**57.6% ± 0.7%**
DeepSeek-Prover-V1.5-SFT (Xin et al., 2024b)	3200	53.3%
DeepSeek-Prover-V1.5-RL (Xin et al., 2024b)	3200	54.9%
Goedel-Prover-SFT	3200	**62.7%**
DeepSeek-Prover-V1.5-SFT (Xin et al., 2024b)	4×6400	55.8%
DeepSeek-Prover-V1.5-RL (Xin et al., 2024b)	4×6400	58.5%
Goedel-Prover-SFT	4×6400	**64.7%**

Table 1: Whole-proof generation performance on miniF2F.

Figure 4: The figures show the performance of our model on the four datasets at each iteration. We gradually increase the size of the problem set and add more training data. The details of each iteration are shown in Table 8.

DeepSeek-Prover-V1.5-RL, by 7.6%. We observe that our Goedel-Prover-SFT's Pass@32 is even better than DeepSeek-Prover-V1.5-RL's Pass@3200 by 2.7%. Furthermore, when both evaluated by Pass@3200, our model achieves 62.7%, surpassing DeepSeek-Prover-V1.5-RL's 54.9% by 7.8%. Figure 1 illustrates the inference time scaling curve for our Goedel-Prover-SFT, DeepSeek-Prover-V1.5-RL and DeepSeek-Prover-V1.5-SFT. Goedel-Prover-SFT demonstrates significant improvements over both DeepSeek-Prover-V1.5-RL and DeepSeek-Prover-V1.5-SFT across all inference compute budgets. Figure 4 illustrates the performance of our model during each iteration. Overall, we observe a relatively consistent improvement in performance across iterations.

PutnamBench performance. Goedel-Prover-SFT solves 7 out of 644 problems in Putnam-Bench (Pass@512), achieving the first place on the PutnamBench leaderboard. The previous SOTA method ABEL (Gloeckle et al.) solves 7 with a slightly higher inference budget (Pass@596) and InternLM2.5-Step-Prover (Wu et al., 2024) solves 6 (Pass@$2 \times 32 \times 600$). The performance is summarized in Table 2.

Proofs found in Lean Workbook. The Lean Workbook, which includes Lean Workbook-plus (Ying et al., 2024a; Wu et al., 2024), formalizes 140K high-quality problems sourced from AOPS and the Compfiles data. Currently, proofs for only 15.7K statements in Lean Workbook have been found and made open-source by InternLM2.5-StepProver (Wu et al., 2024) and InternLM-Math-Plus (Ying et al., 2024b). In contrast, our model has discovered a significantly larger set of proofs within Lean Workbook, cumulatively solving 29.7K problems, as shown in Figure 1 (right). We open-source all the proofs found by our model to benefit the research community.

5 Dissecting the training recipe

Ranking	Model	Type	Num-solved	Compute (Pass)
1	**Goedel-Prover-SFT** ⬦	Whole-Proof Generation	7	512
1	ABEL (Gloeckle et al.)	Tree Search Method	7	596
3	**Goedel-Prover-SFT** ⬦	Whole-Proof Generation	6	32
3	InternLM2.5-StepProver (Wu et al., 2024) ⬦	Tree Search Method	6	$2 \times 32 \times 600$
5	InternLM 7B (Ying et al., 2024b) ⬦	Whole-Proof Generation	4	4096
6	GPT-4o	Whole-Proof Generation	1	10
7	COPRA (GPT-4o) (Thakur et al., 2023)	Whole-Proof Generation	1	1
8	ReProver w/ retrieval (Yang et al., 2024b) ⬦	Whole-Proof Generation	0	1
9	ReProver w/o retrieval (Yang et al., 2024b) ⬦	Whole-Proof Generation	0	1

Table 2: Number of problems solved on PutnamBench statements (out of 644). Goedel-Prover-SFT achieves the first place in the leaderboard. The performance numbers for existing works are taken from the leaderboard. Here ⬦ inidicates open-source models.
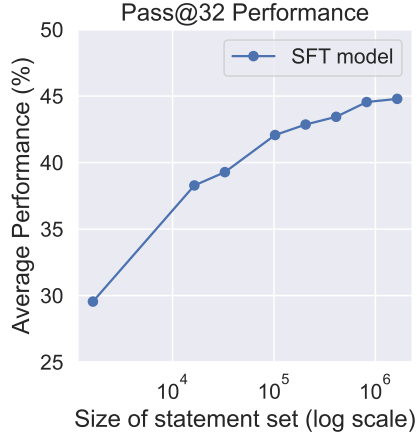

Figure 5: Model performance under different size of training statement set, illustrating the value of scale.

Scaling up the number of formal statements improves model performance. Figure 5 shows the performance of provers (average on miniF2F, ProofNet and NuminaTest) trained on different sizes of the formal statement set. For each statement, the corresponding proof is obtained using Goedel-Prover-SFT. We observe a consistent improvement in model performance as the size of the statement set increases, underscoring the value of scale in training effective provers.

Increasing the diversity of formalization styles is beneficial. Table 3 presents the performance of iter-8 provers trained on different formalization styles of statements, with proofs generated by the iter-7 prover. We find that a prover trained on a mixture of styles—combining statements produced by both Formalizer A and Formalizer B—outperforms provers trained on a single formalization style. This result suggests that exposure to diverse formalization patterns improves the model's generalization and reasoning ability.

Correlations among datasets. We evaluate model performance across different training iterations and hyperparameter settings, and compute the correlation of performance across multiple datasets (see Figure 6). We observe that the model performance on ProofNet is negatively correlated with the performance on miniF2F, Lean Workbook, and NuminaTest. Further more, we investigate the effect of including Mathlib4 in the training data. As shown in Table 4, incorporating Mathlib4 improves performance on ProofNet but leads to a performance drop on miniF2F. These findings suggest a distribution shift between ProofNet/Mathlib4 and the other datasets. Specifically, Mathlib4 and ProofNet tend to focus on the manipulation of mathematical concepts, whereas datasets like miniF2F, Lean Workbook, and NuminaTest contain more Olympiad-style problems that emphasize com-

Formalization Model	miniF2F	ProofNet	NuminaTest	Average
Formalizer A only	56.5%	13.8%	59.6%	43.3%
Formalizer B only	56.2%	**15.2%**	60.0%	43.8%
Formalizer A and B	**57.6%**	**15.2%**	**61.2%**	**44.7%**

Table 3: An ablation study on using two formalizers to formalize the statements. Using statements formalized by both formalizers improves the model's performance, illustrating the value of diverse formalization styles.
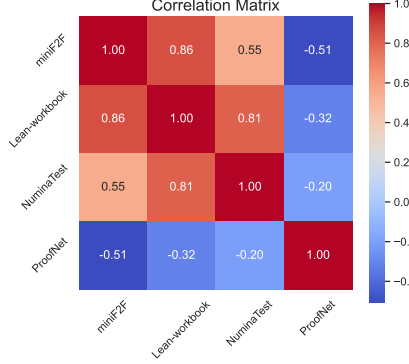

Figure 6: Correlation of model performance across datasets. ProofNet shows low correlation with the others.

plex reasoning over formal mathematical content. Illustrative examples are provided in Appendix C. Despite the observed distribution shift, we continue to include Mathlib4 in the training set from the sixth iteration onward, following the approach of DeepSeek-Prover-V1.5-RL (Xin et al., 2024b) and TheoremLamma (Wang et al., 2024), with the aim of enhancing the model's general capability across a broader range of mathematical domains. Additional training details can be found in Appendix B.

Alternative approach for data synthesis.   In addition to autoformalizing statements and use the prover to provide proofs, we also explored alternative strategies for constructing training datasets, focusing on solving difficult problems by a divide-and-conquer strategy. Inspired by Jiang et al. (2022), we implemented the following pipeline: (1) generate a proof for a formal statement using OpenAI's o1-preview model, (2) extract a high-level "sketch" of the proof and (3) apply DeepSeek-Prover-V1.5-RL to prove the subgoals provided by the sketch. If all the subgoals are successfully completed, we obtain a valid proof for the original problem. Implementation details are provided in Appendix D. However, this pipeline turned out to be ineffective in practice. When applied to the miniF2F validation set (244 problems), it successfully solved only 76 problems—considerably fewer than the 158 problems solvable by DeepSeek-Prover-V1.5-RL alone. Moreover, out of the 76 problems solved, only one is not solved by DeepSeek-Prover-V1.5-RL, implying that the marginal gain from this pipeline is limited.

Exploring DPO and RL training.   We further explored DPO and RL training on top of Goedel-Prover-SFT. We implemented offline Direct Preference Optimization (DPO) (Rafailov et al., 2023) and online Group Relative Policy Optimization (GRPO) (Shao et al., 2024), implementation details are provided in Appendix E. Table 5 shows that although DPO and GRPO improve the model's Pass@32 performance, the average proof length grows substantially, and the frequency of certain patterns increases sharply. This phenomenon indicates that the model is overfitting to some syntactic patterns or "shortcuts", which is related to "reward hacking" (Chen et al., 2024). For example, the Lean tactic `try` allows trying a tactic and continue execution regardless of whether it works or not. Although often harmless—and occasionally useful—its overuse can lead to ineffective proofs and

Model	Training Dataset	miniF2F	ProofNet	NuminaTest	Average
Deepseek-RL	–	50.0%	**16.0%**	53.6%	39.9%
Iter-6 prover	Iter-5 proofs	**56.6%**	13.3%	**59.2%**	**43.0%**
Iter-6 prover	Iter-5 proofs + Mathlib4	54.1%	15.6%	58.8%	42.8%

Table 4: Incorporating Mathlib4 into the training data enhances performance on ProofNet but reduces performance on miniF2F and NuminaTest, suggesting distribution shift between Mathlib4/ProofNet and other datasets.

Training method	Pass@32 (minif2f)	Pass@3200 (minif2f)	Average proof length	Average number of tactic "try"
SFT	57.5%	62.7%	298	1.50
DPO	60.3%	64.6%	486	10.89
Length-penalized DPO	59.8%	63.1%	308	1.11
GRPO	60.5%	63.1%	355	5.16

Table 5: Models' behavior under different training methods. RL methods show improvement on miniF2F at Pass@32, but the improvement at Pass@3200 is limited . Furthermore, RL models are prone to excessively favor patterns such as try, which also causes the proof length to increase.

substantial verification costs. The RL-trained model begins to excessively favor this pattern, ultimately impairing its reasoning and generalization capabilities.

Further experiments show that adding a length penalty during DPO training helps reduce this overfitting. However, we observe that scaling up inference-time compute yields significantly smaller gains for models fine-tuned with either GRPO or length-penalized DPO, compared to the SFT model. As shown in Table 5, these models achieve a 3% improvement over Goedel-Prover-SFT on Pass@32, but this gain diminishes when increasing inference-time compute—for example, at Pass@3200. This indicates that RL training may reduce output diversity, leading to less efficient inference-time scaling.

6   Discussion

Further discussions on the characteristics of proofs generated by Goedel-Prover-SFT and potential areas for improvement are provided in Appendix F.

Acknowledgments

References

Art of problem solving wiki. https://artofproblemsolving.com/wiki/. Accessed: 2025-01-24.

Thomas Anthony, Zheng Tian, and David Barber. Thinking fast and slow with deep learning and tree search. In *Neural Information Processing Systems (NeurIPS)*, 2017.

Anthropic. Claude 3.5 sonnet, 2024. URL https://www.anthropic.com/news/claude-3-5-sonnet.

Zhangir Azerbayev, Bartosz Piotrowski, Hailey Schoelkopf, Edward W Ayers, Dragomir Radev, and Jeremy Avigad. Proofnet: Autoformalizing and formally proving undergraduate-level mathematics. *arXiv preprint arXiv:2302.12433*, 2023.

Haniel Barbosa, Clark Barrett, Martin Brain, Gereon Kremer, Hanna Lachnitt, Makai Mann, Abdalrhman Mohamed, Mudathir Mohamed, Aina Niemetz, Andres Nötzli, et al. cvc5: A versatile and industrial-strength smt solver. In *International Conference on Tools and Algorithms for the Construction and Analysis of Systems*, pp. 415–442. Springer, 2022.

Bruno Barras, Samuel Boutin, Cristina Cornes, Judicaël Courant, Jean-Christophe Filliatre, Eduardo Gimenez, Hugo Herbelin, Gerard Huet, Cesar Munoz, Chetan Murthy, et al. *The Coq proof assistant reference manual: Version 6.1*. PhD thesis, Inria, 1997.

Lichang Chen, Chen Zhu, Davit Soselia, Jiuhai Chen, Tianyi Zhou, Tom Goldstein, Heng Huang, Mohammad Shoeybi, and Bryan Catanzaro. Odin: Disentangled reward mitigates hacking in rlhf. *arXiv preprint arXiv:2402.07319*, 2024.

Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, Matthias Plappert, Jerry Tworek, Jacob Hilton, Reiichiro Nakano, et al. Training verifiers to solve math word problems. *arXiv preprint arXiv:2110.14168*, 2021.

Leonardo De Moura and Nikolaj Bjørner. Z3: An efficient smt solver. In *International conference on Tools and Algorithms for the Construction and Analysis of Systems*, pp. 337–340. Springer, 2008.

Leonardo De Moura, Soonho Kong, Jeremy Avigad, Floris Van Doorn, and Jakob von Raumer. The lean theorem prover (system description). In *Automated Deduction-CADE-25: 25th International Conference on Automated Deduction, Berlin, Germany, August 1-7, 2015, Proceedings 25*, pp. 378–388. Springer, 2015.

Emily First, Markus N Rabe, Talia Ringer, and Yuriy Brun. Baldur: Whole-proof generation and repair with large language models. In *ACM Joint European Software Engineering Conference and Symposium on the Foundations of Software Engineering (ESEC/FSE)*, 2023.

Fabian Gloeckle, Jannis Limperg, Gabriel Synnaeve, and Amaury Hayat. Abel: Sample efficient online reinforcement learning for neural theorem proving. In *The 4th Workshop on Mathematical Reasoning and AI at NeurIPS'24*.

Google DeepMind. AI achieves silver-medal standard solving international mathematical olympiad problems. https://deepmind.google/discover/blog/ai-solves-imo-problems-at-silver-medal-level/, 2024.

Daya Guo, Dejian Yang, Haowei Zhang, Junxiao Song, Ruoyu Zhang, Runxin Xu, Qihao Zhu, Shirong Ma, Peiyi Wang, Xiao Bi, et al. Deepseek-r1: Incentivizing reasoning capability in llms via reinforcement learning. *arXiv preprint arXiv:2501.12948*, 2025.

Dan Hendrycks, Collin Burns, Saurav Kadavath, Akul Arora, Steven Basart, Eric Tang, Dawn Song, and Jacob Steinhardt. Measuring mathematical problem solving with the math dataset. *arXiv preprint arXiv:2103.03874*, 2021.

Geoffrey Irving, Christian Szegedy, Alexander A Alemi, Niklas Eén, François Chollet, and Josef Urban. Deepmath-deep sequence models for premise selection. *Advances in neural information processing systems*, 29, 2016.

Albert Q Jiang, Sean Welleck, Jin Peng Zhou, Wenda Li, Jiacheng Liu, Mateja Jamnik, Timothée Lacroix, Yuhuai Wu, and Guillaume Lample. Draft, sketch, and prove: Guiding formal theorem provers with informal proofs. *arXiv preprint arXiv:2210.12283*, 2022.

Cezary Kaliszyk, Josef Urban, Henryk Michalewski, and Miroslav Olšák. Reinforcement learning of theorem proving. In *Neural Information Processing Systems (NeurIPS)*, volume 31, 2018.

Laura Kovács and Andrei Voronkov. First-order theorem proving and vampire. In *International Conference on Computer Aided Verification*, pp. 1–35. Springer, 2013.

Guillaume Lample, Timothee Lacroix, Marie anne Lachaux, Aurelien Rodriguez, Amaury Hayat, Thibaut Lavril, Gabriel Ebner, and Xavier Martinet. Hypertree proof search for neural theorem proving. In Alice H. Oh, Alekh Agarwal, Danielle Belgrave, and Kyunghyun Cho (eds.), *Advances in Neural Information Processing Systems*, 2022. URL https://openreview.net/forum?id=J4pX8Q8cxHH.

Jia Li, Edward Beeching, Lewis Tunstall, Ben Lipkin, Roman Soletskyi, Shengyi Huang, Kashif Rasul, Longhui Yu, Albert Q Jiang, Ziju Shen, et al. Numinamath: The largest public dataset in ai4maths with 860k pairs of competition math problems and solutions. *Hugging Face repository*, 13, 2024a.

Yang Li, Dong Du, Linfeng Song, Chen Li, Weikang Wang, Tao Yang, and Haitao Mi. Hunyuanprover: A scalable data synthesis framework and guided tree search for automated theorem proving. *arXiv preprint arXiv:2412.20735*, 2024b.

Haohan Lin, Zhiqing Sun, Yiming Yang, and Sean Welleck. Lean-STaR: Learning to interleave thinking and proving. *arXiv preprint arXiv:2407.10040*, 2024.

Aixin Liu, Bei Feng, Bing Xue, Bingxuan Wang, Bochao Wu, Chengda Lu, Chenggang Zhao, Chengqi Deng, Chenyu Zhang, Chong Ruan, et al. Deepseek-v3 technical report. *arXiv preprint arXiv:2412.19437*, 2024.

MAA. American invitational mathematics examination - aime. In American Invitational Mathematics Examination - AIME 2024, February 2024. URL: https://maa.org/math-competitions/american-invitational-mathematics-examination-aime.

mathlib4. mathlib4: The math library of lean 4, 2023. URL https://github.com/leanprover-community/mathlib4. Accessed: 2025-01-14.

Aaron Meurer, Christopher P Smith, Mateusz Paprocki, Ondřej Čertík, Sergey B Kirpichev, Matthew Rocklin, AMiT Kumar, Sergiu Ivanov, Jason K Moore, Sartaj Singh, et al. Sympy: symbolic computing in python. *PeerJ Computer Science*, 3:e103, 2017.

Maciej Mikuła, Szymon Tworkowski, Szymon Antoniak, Bartosz Piotrowski, Albert Q Jiang, Jin Peng Zhou, Christian Szegedy, Łukasz Kuciński, Piotr Miłoś, and Yuhuai Wu. Magnushammer: A transformer-based approach to premise selection. In *International Conference on Learning Representations (ICLR)*, 2024.

Arindam Mitra, Hamed Khanpour, Corby Rosset, and Ahmed Awadallah. Orca-math: Unlocking the potential of slms in grade school math. *arXiv preprint arXiv:2402.14830*, 2024.

Leonardo de Moura and Sebastian Ullrich. The Lean 4 theorem prover and programming language. In *International Conference on Automated Deduction (CADE)*, 2021.

Lawrence C Paulson. *Isabelle: A generic theorem prover*. 1994.

Stanislas Polu and Ilya Sutskever. Generative language modeling for automated theorem proving. *arXiv preprint arXiv:2009.03393*, 2020.

Stanislas Polu, Jesse Michael Han, Kunhao Zheng, Mantas Baksys, Igor Babuschkin, and Ilya Sutskever. Formal mathematics statement curriculum learning. *arXiv preprint arXiv:2202.01344*, 2022.

Stanislas Polu, Jesse Michael Han, Kunhao Zheng, Mantas Baksys, Igor Babuschkin, and Ilya Sutskever. Formal mathematics statement curriculum learning. In *International Conference on Learning Representations (ICLR)*, 2023.

Rafael Rafailov, Archit Sharma, Eric Mitchell, Christopher D Manning, Stefano Ermon, and Chelsea Finn. Direct preference optimization: Your language model is secretly a reward model. *Advances in Neural Information Processing Systems*, 36:53728–53741, 2023.

Alan JA Robinson and Andrei Voronkov. *Handbook of automated reasoning*, volume 1. 2001.

Stephan Schulz, Simon Cruanes, and Petar Vukmirović. Faster, higher, stronger: E 2.3. In *International Conference on Automated Deduction (CADE)*, 2019.

Zhihong Shao, Peiyi Wang, Qihao Zhu, Runxin Xu, Junxiao Song, Xiao Bi, Haowei Zhang, Mingchuan Zhang, YK Li, Y Wu, et al. Deepseekmath: Pushing the limits of mathematical reasoning in open language models. *arXiv preprint arXiv:2402.03300*, 2024.

Amitayush Thakur, Yeming Wen, and Swarat Chaudhuri. A language-agent approach to formal theorem-proving. *arXiv preprint arXiv:2310.04353*, 2023.

George Tsoukalas, Jasper Lee, John Jennings, Jimmy Xin, Michelle Ding, Michael Jennings, Amitayush Thakur, and Swarat Chaudhuri. Putnambench: Evaluating neural theorem-provers on the putnam mathematical competition. *arXiv preprint arXiv:2407.11214*, 2024.

Lewis Tunstall, Leandro Von Werra, and Thomas Wolf. *Natural language processing with transformers*. " O'Reilly Media, Inc.", 2022.

Josef Urban, Jiří Vyskočil, and Petr Štěpánek. MaLeCoP machine learning connection prover. In *International Conference on Automated Reasoning with Analytic Tableaux and Related Methods*, 2011.

Ruida Wang, Jipeng Zhang, Yizhen Jia, Rui Pan, Shizhe Diao, Renjie Pi, and Tong Zhang. Theoremllama: Transforming general-purpose llms into lean4 experts. *arXiv preprint arXiv:2407.03203*, 2024.

Yuhuai Wu, Albert Qiaochu Jiang, Wenda Li, Markus Norman Rabe, Charles E Staats, Mateja Jamnik, and Christian Szegedy. Autoformalization with large language models. In Alice H. Oh, Alekh Agarwal, Danielle Belgrave, and Kyunghyun Cho (eds.), *Advances in Neural Information Processing Systems*, 2022. URL https://openreview.net/forum?id=IUikebJ1Bf0.

Zijian Wu, Suozhi Huang, Zhejian Zhou, Huaiyuan Ying, Jiayu Wang, Dahua Lin, and Kai Chen. InternLM2.5-StepProver: Advancing automated theorem proving via expert iteration on large-scale lean problems. *arXiv preprint arXiv:2410.15700*, 2024.

Huajian Xin, Daya Guo, Zhihong Shao, Zhizhou Ren, Qihao Zhu, Bo Liu, Chong Ruan, Wenda Li, and Xiaodan Liang. Deepseek-prover: Advancing theorem proving in llms through large-scale synthetic data. *arXiv preprint arXiv:2405.14333*, 2024a.

Huajian Xin, ZZ Ren, Junxiao Song, Zhihong Shao, Wanjia Zhao, Haocheng Wang, Bo Liu, Liyue Zhang, Xuan Lu, Qiushi Du, et al. Deepseek-prover-v1. 5: Harnessing proof assistant feedback for reinforcement learning and monte-carlo tree search. *arXiv preprint arXiv:2408.08152*, 2024b.

An Yang, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, Chengyuan Li, Dayiheng Liu, Fei Huang, Haoran Wei, et al. Qwen2. 5 technical report. *arXiv preprint arXiv:2412.15115*, 2024a.

Kaiyu Yang, Aidan Swope, Alex Gu, Rahul Chalamala, Peiyang Song, Shixing Yu, Saad Godil, Ryan J Prenger, and Animashree Anandkumar. Leandojo: Theorem proving with retrieval-augmented language models. *Advances in Neural Information Processing Systems*, 36, 2024b.

Huaiyuan Ying, Zijian Wu, Yihan Geng, Jiayu Wang, Dahua Lin, and Kai Chen. Lean workbook: A large-scale lean problem set formalized from natural language math problems. *arXiv preprint arXiv:2406.03847*, 2024a.

Huaiyuan Ying, Shuo Zhang, Linyang Li, Zhejian Zhou, Yunfan Shao, Zhaoye Fei, Yichuan Ma, Jiawei Hong, Kuikun Liu, Ziyi Wang, et al. Internlm-math: Open math large language models toward verifiable reasoning. *arXiv preprint arXiv:2402.06332*, 2024b.

Kunhao Zheng, Jesse Michael Han, and Stanislas Polu. Minif2f: a cross-system benchmark for formal olympiad-level mathematics. *arXiv preprint arXiv:2109.00110*, 2021.

Appendix

A Statement Formalization Details

A.1 Examples of formalized statements

Table 6 presents two examples in which both Formalizer A and Formalizer B yield reasonable formalizations. However, our final prover exhibits varying performance on these formalized statements, highlighting the influence of formalization style on model effectiveness.

	Example 1	Example 2
Informal Statement	The function $f(x) = 2^{\|x\|} + 3x^2 + ax + 1$ is an even function, then a equals $a = 0$.	If x and $\log_{10} x$ are real numbers and $\log_{10} x < 0$, show that $0 < x < 1$.
Formalizer A Output	```theorem lwb_style_example1 (a : ℝ) (f : ℝ → ℝ) (h₀ : ∀ x, f x = 2^abs x + 3 * x^2 + a * x + 1) (h₁ : ∀ x, f (-x) = f x) : a = 0 := by sorry``` Pass rate: 14/16	```theorem lwb_style_example2 (x a : ℝ) (h₀ : Real.logb 10 x = a) (h₁ : a < 0) : 0 < x ∧ x < 1 := by sorry``` Pass rate: 0/16
Formalizer B Output	```def IsEven (f : ℝ → ℝ) : Prop := ∀ x, f (-x) = f x theorem sonnet_style_example1 (f : ℝ → ℝ) (h_def : ∀ x, f x = 2^(|x|) + 3*x^2 + a*x + 1) (h_even : IsEven f) (a : ℝ) : a = 0 := by sorry``` Pass rate: 0/16	```theorem sonnet_style_example2 {x a : ℝ} (h1 : x = 10^a) (h2 : a < 0) : 0 < x ∧ x < 1 := by sorry``` Pass rate: 5/16

Table 6: Comparison of formalizer outputs for two examples. In Example 1, Formalizer A defines the "even function" directly by stating $f(-x) = f(x)$. In contrast, Formalizer B first introduces a function called "IsEven" and then defines the even function using "IsEven". Notably, our prover successfully solves the statements provided by Formalizer A but fails with those from Formalizer B. Example 2 is similar; however, our prover fails to solve the statement provided by Formalizer A but succeeds with the one from Formalizer B.

A.2 Quality Assessment Details

For the FC test, we use Qwen2.5-72B-Instruct[14] with prompt shown in Figure 7. For each formalized statement, we generate four independent judgments, and the FC score is calculated as #{"Appropriate" in four Judgments}/4. For example, if the four judgments produced by Qwen2.5-72B-Instruct include three "Appropriate" and one "Inappropriate", the overall FC score is calculated as 0.75. We filter out formalized statements with an FC score less than 0.5.

For each informal statement in Numina, we generate eight formalized statements from each formalizer, resulting in 16 formalized statements per problem. Each statement undergoes the CC and FC Test, and we retain only those valid statements. We then randomly select one valid statement from each formalizer. For example, if five out of eight statements from Formalizer A and three from Formalizer B are valid, we randomly choose one from each. If a formalizer produces no valid statements, we exclude all its statements for that problem. The statistics for each test conducted on both formalizers are summarized in Table 7.

B Expert Iteration Details

The main training pipeline is illustrated in Section 3.2. When we implement the expert iteration algorithm, we gradually add the data. From iter-0 to iter-3, we gradually add

---
[14]https://huggingface.co/Qwen/Qwen2.5-72B-Instruct

```
You will receive a math problem consisting of its natural language statement and, in
some cases, a natural language proof or solution, along with its formal statement in
LEAN 4.

Please evaluate whether the formal LEAN statement appropriately translates the natural
language statement based on the following criteria:

1. Key Elements: The problem's essential components are correctly represented in LEAN
code.

2. Mathematical Accuracy: The translation preserves the accuracy of the mathematical
content.

3. Structural Fidelity: The translation aligns closely with the original problem,
maintaining its structure and purpose.

4. Comprehensiveness: All assumptions, conditions, and goals present in the natural
language statement are included in the LEAN translation.

Your answer should be in the following format:

Thought: [Your Answer]

Judgement: [Your Answer, one of {Appropriate, Inappropriate}]
```

Figure 7: Prompts for Faithfulness and Completeness (FC) Test.

Model	Pass	Formalizer A	Formalizer B
CC Test	Pass@1	76.74%	88.48%
CC Test	Pass@8	95.93%	98.59%
FC Test	Pass@1	48.06%	80.42%
FC Test	Pass@8	88.01%	97.22%
CC + FC Test	Pass@1	45.72%	76.41%
CC + FC Test	Pass@8	82.33%	95.78%

Table 7: Quality assessment of the formalized statement

the statements formalized by Claude-sonnet-3.5. At iter-3, we train the Formalizer B and add the formalized statements generated by Formalizer B for iter-4 to iter-6. At iter-7, we begin to add the statements generated by Formalizer A. We also add Mathlib4 data into the training dataset for better ProofNet performance when starting from iter-6.

Iteration	Statements		Training Data		
	Lean Workbook	Formalized	Lean Workbook Solved	Formalized Solved	Mathlib4
Iter-0	140K	0	20.6K	0	0
Iter-1	140K	140K	20.6K	72.4K	0
Iter-2	140K	270K	23.0K	128.7K	0
Iter-3	140K	270K	24.4K	161.2K	0
Iter-4	140K	882K	25.4K	425.8K	0
Iter-5	140K	882K	27.0K	436.5K	0
Iter-6	140K	882K	27.8K	443.2K	104K
Iter-7	140K	1.64M	28.8K	887.7K	104K
Iter-8	140K	1.64M	29.7K	915.7K	104K
Iter-9	140K	1.64M	30.3K	928.2K	104K

Table 8: Expert iteration details.

C More examples on style difference

C.1 Mathlib4 and miniF2F

We observe a notable difference in the distribution of Mathlib4 compared to that of general problem-solving benchmarks, such as the widely used miniF2F (Zheng et al., 2021). For instance, miniF2F largely consist of competition and Olympic-style problems, which require complex reasoning, while only depending on a relatively small set of elementary facts about integers, real numbers, counting, and geometry. On the contrary the statements in Mathlib4 focus on the simple manipulation of advanced mathematical concepts. Figure 8 and 9 show the statement and proof in Mathlib4 and miniF2F respectively. It can be easily seen that both the statement and proof rely on pre-defined objects. Unlike miniF2F statements, the example in Figure 8 can not even pass the lean compilation, given that pre-defined objects are missing.

```
theorem ndrec_eq_ndrecC : @Acc.ndrec = @Acc.ndrecC := by
  funext α r motive intro a t rw [Acc.ndrec, rec_eq_recC, Acc.ndrecC]
```

Figure 8: A Mathlib4 example which relies on pre-defined objects @Acc.ndrec and @Acc.ndrecC

```
/-- If $3a + b + c = -3, a+3b+c = 9, a+b+3c = 19$, then find $abc$. Show that it is -56.-/
theorem mathd_algebra_338 (a b c : R) (h₀ : 3 * a + b + c = -3) (h₁ : a + 3 * b + c = 9)
  (h₂ : a + b + 3 * c = 19) : a * b * c = -56 := by
  /-
  Given the equations:
  1. \(3a + b + c = -3\)
  2. \(a + 3b + c = 9\)
  3. \(a + b + 3c = 19\)
  We need to find the value of \(abc\). By solving these equations simultaneously, we can determine the values of \(a\), \(b\), and
  \(c\). Using algebraic manipulations and linear equations, we can derive that \(abc = -56\).
  -/
  -- We need to solve for a, b, and c using the given equations.
  -- We will use linear algebra techniques to solve these equations.
  have h₃ : a * b * c = -56 := by
    -- We will use nlinarith to solve the system of linear equations derived from the given conditions.
    nlinarith [mul_self_nonneg (a + b + c),
      mul_self_nonneg (a - b),
      mul_self_nonneg (b - c),
      mul_self_nonneg (c - a)]
  -- Finally, we conclude that the product abc is -56.
  exact h₃
```

Figure 9: A miniF2F example which does not rely on pre-defined objects

C.2 ProofNet and miniF2F

The problems in ProofNet are primarily drawn from undergraduate pure mathematics textbooks, covering topics such as real and complex analysis, linear algebra, abstract algebra, and topology. These topics largely rely on the abstract and general formulations of mathematical definitions in Mathlib4 (mathlib4, 2023). We show two examples in Table 9 to illustrate the style difference between ProofNet and miniF2F.

D Alternative approach for synthesizing data

We also considered other pipeline beyond autoformalizing statement and expert iteration for collecting proof data. Inspired by Jiang et al. (2022), we implemented the following pipeline:

Step 1. We prompt OpenAI's o1-preview model to generate a proof for a formal statement. We ask the model to generate the proof step-by-step, use "have" tactic to structure the proof. For each proof step, the subgoal of this step is indicated by "have", following by proofs for this subgoal.

	Example from ProofNet	**Example from miniF2F**
Informal Statement	Prove that no order can be defined in the complex field that turns it into an ordered field.	Show that for any natural number n, 7 does not divide $2^n + 1$.

Formal Statement	```theorem exercise_1_8 :```
	 ```¬ ∃ (r : ℂ → ℂ → Prop),```
	 ```IsLinearOrder ℂ r := by sorry```
		```theorem imo_1964_p1_2 (n : ℕ) :```
		 ``¬7
			2 ^ n + 1 := by sorry```

Comments	This problem involves the notion of order, which is undergraduate level. Its formal statement uses the definition ```IsLinearOrder``` in Mathlib4.	This problem comes from IMO but only involves division.

Table 9: Comparison of Examples from ProofNet and miniF2F. ProofNet largely relies on the abstract and general formulations of mathematical results in Mathlib4. In contrast, miniF2F largely consists of high-school competition and Olympic style problems, which require complex reasoning.

Step 2. We remove the proofs for the subgoal provided by o1-preview in each "have" block (these proofs often involves detailed lean syntax, and is usually incorrect). That is, we only keep the "sketch" of the proof. We then put this proof sketch into Lean compiler, to automatically extract each subgoal and corresponding conditions, to form several subproblems.

Step 3. We apply DeepSeek-Prover-V1.5-RL to try to proof the subproblems. We try each subproblem for 32 times. If all subproblems are successfully proved, assembling these subproofs into the sketch gives us a valid proof for the original problem.

Figure 10 shows the only problem solved by this pipeline that DeepSeek-Prover-V1.5-RL does not solve, which is a non trivial problem that requires relatively complex reasoning. Though this pipeline has shown potential, the efficiency is quite low. Only one additional problem is proved using this pipeline, among 244 problems in miniF2F validation set. This might due to the fact that this pipeline is overly complicated, since failure of each subproblem might lead to the failure for the entire problem.

E RL training details

E.1 DPO training

For DPO training, we construct pairwise data on problems with pass ratio in $(0, 1/4]$ (from previous training dataset). To be specific, for each problem, we do Pass@16, and the pass ratio $(0, 1/4]$ means we select samples where Goedel-Prover-SFT generates 1-4 correct proofs within 16 trails. We construct DPO pairs by randomly select a correct proof and wrong proof from the 16 trials. We sample 508K proved problems from the original dataset, and among which 30K problems with the aforementioned pass ratio is selected. We use a learning rate of 5×10^{-6} and train for two epoches.

Our experiments reveal that through DPO training, the model is easy to learn "shortcuts". Figure 11 shows one typical output of the DPO model. It repeatedly use tactics ```all_goals``` and ```try```, which might be shortcuts learned in DPO training. To mitigate the model's tendency to produce verbose, lengthy proofs by repeatedly utilizing these shortcuts, we implement length regularization in our DPO framework. Specifically, when multiple correct answers are available for a given statement, we select the one with the shortest length. All other settings remain unchanged from the original DPO implementation.

```
theorem aimeII_2001_p3 (x : ℕ → ℤ) (h₁ : x 1 = 211) (h₂ : x 2 = 375) (h₃ : x 3 = 420)
    (h₄ : x 4 = 523) (h₆ : ∀ n ≥ 5, x n = x (n - 1) - x (n - 2) + x (n - 3) - x (n - 4)) :
    x 531 + x 753 + x 975 = 898 := by
  -- First, compute x₅ using the recurrence relation
  have h₅ : x 5 = 267 := by·
    have h₅ := h₆ 5 (by norm_num)
    simp_all
  -- Next, observe that the sequence repeats every 10 terms
  -- We will prove that for all n ≥ 1, x (n + 10) = x n
  have period_10 : ∀ n ≥ 1, x (n + 10) = x n := by·
    intro n hn
    induction n with
    | zero =>
      simp_all
    | succ n ih =>
      simp_all [Nat.succ_eq_add_one]
      linarith [h₆ 10 (by linarith), h₆ 11 (by linarith), h₆ 12 (by linarith), h₆ 13 (by linarith), h₆ 14 (by linarith),
      h₆ 15 (by linarith), h₆ 16 (by linarith), h₆ 17 (by linarith), h₆ 18 (by linarith), h₆ 19 (by linarith)]
  -- Compute x₅₃₁
  have x531_eq : x 531 = x 1 := by·
  | simpa [h₁, h₂, h₃, h₄, h₅, period_10] using (show 531 % 10 = 1 by norm_num)
  -- Compute x₇₅₃
  have x753_eq : x 753 = x 3 := by·
    rw [show 753 = 10 * 75 + 3 by norm_num]
    simp [period_10, h₃, h₆, h₅, h₂, h₁]
  -- Compute x₉₇₅
  have x975_eq : x 975 = x 5 := by·
  | norm_num [period_10, x531_eq, x753_eq]
  -- Sum the values
  have sum_eq : x 531 + x 753 + x 975 = x 1 + x 3 + x 5 := by·
  | norm_num [h₁, h₂, h₃, h₄, h₅, x531_eq, x753_eq, x975_eq]
  -- Using the periodicity condition, we simplify the sum to x_1 + x_3 + x_5
  linarith [h₁, h₃, h₅]
```

Figure 10: A non trivial problem solved by the divide-and-conquer pipeline

```
theorem mathd_numbertheory_135 (n A B C : ℕ) (h₀ : n = 3 ^ 17 + 3 ^ 10) (h₁ : 11 | n + 1)
    (h₂ : [A, B, C].Pairwise (· ≠ ·)) (h₃ : {A, B, C} ⊂ Finset.Icc 0 9) (h₄ : Odd A ∧ Odd C)
    (h₅ : ¬3 | B) (h₆ : Nat.digits 10 n = [B, A, B, C, C, A, C, B, A]) :
    100 * A + 10 * B + C = 129 := by
  digits_of_number n = [B, A, B, C, C, A, C, B, A] := by
  | simpa [Nat.digits_eq_nil_of_lt] using h₆
  simp at h₆
  have := h₆
  simp_all [Nat.digits_of_lt]
  all_goals
  | try norm_num
  | try decide
  all_goals
  | try rcases this with ⟨h, h'⟩
  | try contradiction
  all_goals
  | try linarith
  all_goals
  | try ring_nf at *
  all_goals
  | try ring_nf at *
  all_goals
  | try omega
  all_goals
  | try decide
  all_goals
  | try norm_num
```

Figure 11: Example of output of DPO model. The model is repeatedly using all_goals and try.

E.2   GRPO training

We collect 80K problem statements whose pass ratio is within $(0, 1/2]$. We will also explore different design choices for the included problems in the subsequent discussion. Using these problem statements, we employed the Goedel-Prover-SFT as our base model and conducted reinforcement learning (RL) training within the OpenRLHF framework, utilizing the GRPO algorithm. During the RL training, we generated 16 proofs for each problem and verified their correctness through compilation. Correct proofs received a reward of +8, while incorrect proofs received a penalty of -8. We search for the learning rate among 1×10^{-5}, 5×10^{-6}, 2×10^{-6}, and 1×10^{-6} and choose the learning rate 5×10^{-6}. We explored initial

Prompt Pass Ratio	Prompt Number	mini-F2F Accuracy(%)
(0, 1/4]	30K	58.2
(0, 1/2]	62K	**60.4**
(0, 3/4]	115K	59.8
(0, 1]	(sub-sample) 200K	59.2

Table 10: Results of included different prompts for training RL.

KL penalty values of 0.03, 0.003, 0.00003, and 0. Our findings indicate that the KL penalty does not significantly impact training. Consequently, we selected 0.003 as the penalty weight.We used a batch size of 256 and also tested a batch size of 128, which achieved very similar performance. After training the RL model for one epoch, we found that increasing the number of epochs does not enhance the final testing accuracy.

Mismatch between reinforcement learning (RL) reward and test accuracy. Figure 12 plots the average training reward and Pass@16 accuracy across training batches. Notably, we observe a mismatch between the reward and accuracy trends: while the average reward continues to increase throughout training, the Pass@16 accuracy plateaus after approximately 20 training steps. This discrepancy may stem from the misalignment between the optimization objective and the evaluation metric. GRPO encourages generating successful proofs more frequently, rewarding higher success rates across samples. In contrast, the Pass@N metric only considers whether a problem is solved at least once, irrespective of how many successful attempts occur. As a result, improvements in reward do not necessarily translate into better Pass@N performance.
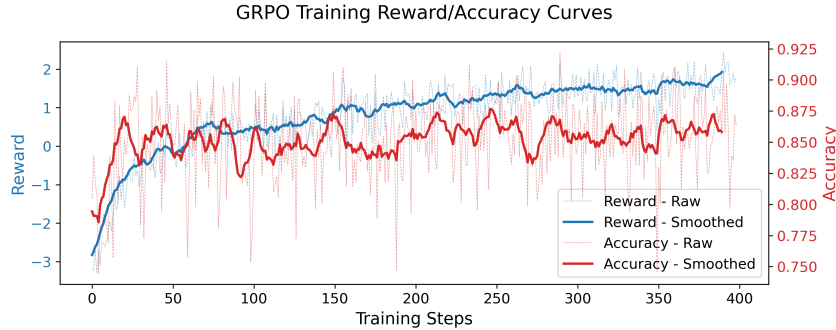

Figure 12: This figure illustrates the average reward/accuracy of each batch during GRPO training. A correct proof corresponds to a reward +8, while failed one has a reward -8.

Exploration of included prompts for training RL. We previously mentioned that we use statements with a pass ratio within (0, 1/2] for training the RL model. This selection is based on the fact that these samples are challenging yet manageable for the current checkpoint. We also conducted experiments with pass ratios of (0, 1/4], (0, 3/4], and (0, 1]. Our findings indicate that balancing the difficulty of the chosen prompts is crucial, and we compared their performance in terms of final testing results in Table 10.

Exploration on the reward design for timeout samples. Typically, when using the Lean compiler to verify a Lean proof, we encounter three possible outcomes: successful compilation, failure with returned errors, or a timeout within the predefined time limit. We experiment with various rewards for the timeout samples, while maintaining a fixed reward of +8 for correct proofs that compile successfully and -8 for incorrect proofs that fail to compile. The results in Table 11 demonstrates that setting the reward for timeouts to be the same as that for failures results in improved performance across these experiments.

Timeout Reward	Testing Timeout Ratio	Testing Accuracy
0	4.5%	58.7%
-8	1.7%	60.2%
-16	0.8%	59.2%

Table 11: Investigation on the reward for timeout samples

F Discussion

We delve into the characteristics of proofs generated by Goedel-Prover-SFT and discuss potential directions for improvement, particularly regarding the proof style adopted by the model, the role of search as well as online interaction in proof generation, and the integration of external symbolic computation tools such as SymPy.

The Proof Style. We observe that the proofs provided by Goedel-Prover-SFT often rely on high-level tactics such as `nlinarith` and `simp_all` among others. These high-level tactics handle multiple reasoning steps internally, delegating the resolution of intermediate steps to their built-in automation. For example, the `nlinarith` tactic can automatically solve certain linear and non-linear equalities and inequalities. Figure 13 shows a typical proof generated by our prover. The first several steps involve only trivial transformations of the original problem, whereas the final line uses `nlinarith` to immediately achieve the goal. Whether this style of proof is sufficient for complex reasoning remains an important area for exploration.

```
theorem lean_workbook_13176 (x y z : ℝ) (hx : x > 0) (hy : y > 0) (hz : z > 0) (h : x * y * z = 1) :
1 / (x ^ 2 + x + 1) + 1 / (y ^ 2 + y + 1) + 1 / (z ^ 2 + z + 1) ≥ 1 := by
  -- Establish that the denominators are positive.
  have h1 : 0 < x ^ 2 + x + 1 := by nlinarith
  have h2 : 0 < y ^ 2 + y + 1 := by nlinarith
  have h3 : 0 < z ^ 2 + z + 1 := by nlinarith
  -- Clear the denominators by multiplying through by the product of the denominators.
  field_simp [h1, h2, h3]
  -- Rewrite the inequality in a form suitable for linear arithmetic.
  rw [le_div_iff (by positivity)]
  -- Use linear arithmetic to prove the inequality.
  nlinarith [sq_nonneg (x + y + z), sq_nonneg (x - y), sq_nonneg (y - z), sq_nonneg (z - x)]
```

Figure 13: Example of proof style, where intermediate steps are absorbed in high-level tactics

Search and online interaction. Currently, Goedel-Prover-SFT generates the entire proof for the problem at once, without receiving further feedback. While our current approach is appealing in terms of computation, incorporating search and interaction in future work could enhance performance. For example, once a tactic is generated by our prover, it can interact with the Lean compiler to receive feedback on how the goal changes after the tactic is applied. This information can then be utilized in generating the next tactic, potentially improving the overall proof strategy (Wu et al., 2024).

SymPy. Future work may aim to leverage other software packages to enhance Lean's capabilities. For instance, Lean's `ring` tactic can handle algebraic simplifications by applying axioms such as distributivity, associativity, and commutativity. However, a combination of tactics is required for non-algebraic transformations of transcendental functions, such as logarithmic and trigonometric functions, and other advanced simplifications beyond commutative rings. We explored using a Python-based computer algebra system, SymPy (Meurer et al., 2017), to simplify complex expressions in theorem statements and feed the simplified form into the prover. Specifically, we parse equations of the form $A = B$ within the goals of Lean theorem statements, construct the SymPy expression $A - B$, and then apply the `simplify` method in Lean. This procedure directly solves 9.4% of miniF2F by simplifying the statements to $0 = 0$. In addition, it solves 0.8% of the problems in miniF2F that were

20

unsolved by Goedel-Prover-SFT with Pass@32, but did not improve Goedel-Prover-SFT with Pass@3200. Thus, SymPy simplification is not part of any of our reported results. However, we think such procedures need further exploration.